



Quantum[®]Leaps
innovating embedded systems



Application Note

QP[™] and ARM Cortex-M with GNU

Document Revision I
January 2012



Copyright © Quantum Leaps, LLC

www.quantum-leaps.com
www.state-machine.com



Table of Contents

1 Introduction.....	1
1.1 About QP™.....	2
1.2 About the ARM Cortex-M Port.....	3
1.3 Cortex Microcontroller Software Interface Standard (CMSIS).....	3
1.4 Licensing QP.....	3
2 Directories and Files.....	4
2.1 Building the QP Libraries.....	6
2.2 Building and Debugging the Examples.....	6
3 The Vanilla Port.....	8
3.1 The qep_port.h Header File.....	8
3.2 The QF Port Header File.....	8
3.3 Handling Interrupts in the Non-Preemptive Vanilla Kernel.....	9
3.4 Idle Loop Customization in the “Vanilla” Port.....	11
4 The QK Port.....	13
4.1 Single-Stack, Preemptive Multitasking on ARM Cortex-M.....	13
4.2 The QK Port Header File.....	17
4.3 QK Platform-Specific Code for ARM Cortex-M.....	19
4.4 Setting up and Starting Interrupts in QF_onStartup().....	23
4.5 Writing ISRs for QK.....	24
4.6 QK Idle Processing Customization in QK_onIdle().....	24
4.7 Testing QK Preemption Scenarios.....	26
5 QS Software Tracing Instrumentation.....	29
5.1 QS Time Stamp Callback QS_onGetTime().....	30
5.2 QS Trace Output in QF_onIdle()/QK_onIdle().....	31
5.3 Invoking the QSpy Host Application.....	32
6 Related Documents and References.....	33
7 Contact Information.....	34




1 Introduction

This Application Note describes how to use the QP™ state machine framework with the ARM Cortex-M processors (Cortex M0/M1/M3/M4 based on the ARMv7-M architecture [ARMv7-M]). This document describes the following two main implementation options:

1. The cooperative “Vanilla” kernel available in the QF real-time framework; and
2. The preemptive run-to-completion QK kernel.

To focus the discussion, this Application Note uses Code Sourcery CodeBench (version **2011.07-60**), and the EK-LM3S811 board from Texas Instruments, as shown in [Figure 1](#). However, the source code for the QP port described here is generic for all ARM Cortex-M devices and runs without modifications on all ARM Cortex-M cores. The port assumes QP version **4.3.00** or higher.

Figure 1: The EK-LM3S811 evaluation board Rev C used to test the ARM Cortex-M port.



1.1 About QP™

QP™ is a family of very lightweight, open source, state machine-based frameworks for developing event-driven applications. QP enables building well-structured embedded applications as a set of concurrently executing hierarchical state machines (UML statecharts) directly in C or C++ **without big tools**. QP is described in great detail in the book “*Practical UML Statecharts in C/C++, Second Edition: Event-Driven Programming for Embedded Systems*” [PSiCC2] (Newnes, 2008).

As shown in Figure 2, QP consists of a universal UML-compliant event processor (QEP), a portable real-time framework (QF), a tiny run-to-completion kernel (QK), and software tracing instrumentation (QS). Current versions of QP include: QP/C™ and QP/C++™, which require about 4KB of code and a few hundred bytes of RAM, and the ultra-lightweight QP-nano, which requires only 1-2KB of code and just several bytes of RAM.


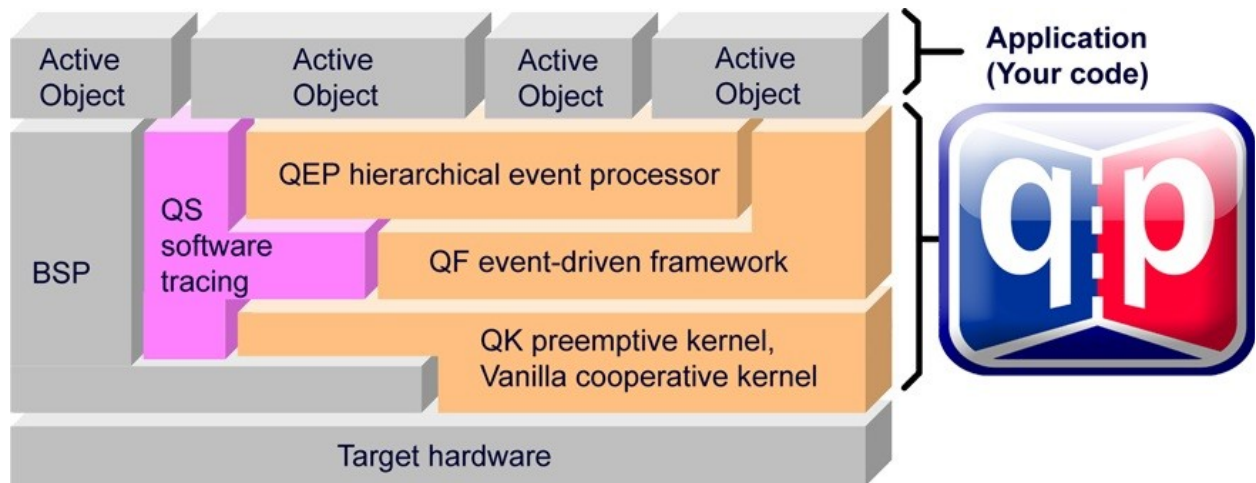


Figure 2: QP components and their relationship with the target hardware, board support package (BSP), and the application



QP can work with or without a traditional RTOS or OS. In the simplest configuration, QP can completely **replace** a traditional RTOS. QP includes a simple non-preemptive scheduler and a fully preemptive kernel (QK). QK is smaller and faster than most traditional preemptive kernels or RTOS, yet offers fully deterministic, preemptive execution of embedded applications. QP can manage up to 63 concurrently executing tasks structured as state machines (called active objects in UML).

QP/C and QP/C++ can also work with a traditional OS/RTOS to take advantage of existing device drivers, communication stacks, and other middleware. QP has been ported to Linux/BSD, Windows, VxWorks, ThreadX, uC/OS-II, FreeRTOS.org, and other popular OS/RTOS.

1.2 About the ARM Cortex-M Port

In contrast to the traditional ARM7/ARM9 cores, ARM Cortex-M cores contain such standard components as the Nested Vectored Interrupt Controller (NVIC) and the System Timer (SysTick). With the provision of these standard components, it is now possible to provide fully portable system-level software for ARM Cortex-M. Therefore, this QP port to ARM Cortex-M can be much more complete than a port to the traditional ARM7/ARM9 and the software is guaranteed to work on any ARM Cortex-M silicon.

The non preemptive cooperative kernel implementation is very simple on ARM Cortex-M, perhaps simpler than any other processor, mainly because Interrupt Service Routines (ISRs) are regular C-functions on ARM Cortex-M.

However, when it comes to handling preemptive multitasking, ARM Cortex-M is a unique processor unlike any other. The ARM Cortex-M hardware has been designed with traditional blocking real-time kernels in mind, and implementing a simple run-to-completion preemptive kernel (such as the QK preemptive kernel described in Chapter 10 in [PSiCC2]) is a little more involved. Section 4 of this application note describes in detail the QK implementation on ARM Cortex-M.

NOTE: This Application Note pertains both to C and C++ versions of the QP™ state machine frameworks. Most of the code listings in this document refer to the C version. Occasionally the C code is followed by the equivalent C++ implementation to show the C++ differences whenever such differences become important.

1.3 Cortex Microcontroller Software Interface Standard (CMSIS)

The ARM Cortex-M examples provided with this Application Note are compliant with the Cortex Microcontroller Software Interface Standard (CMSIS).



1.4 Licensing QP

The **Generally Available (GA)** distributions of QP available for download from the www.state-machine.com/downloads website are offered under the same licensing options as the QP baseline code. These available licenses are:

- The GNU General Public License version 2 (GPL) as published by the Free Software Foundation and appearing in the file `GPL.TXT` included in the packaging of every Quantum Leaps software distribution. The GPL *open source* license allows you to use the software at no charge under the condition that if you redistribute the original software or applications derived from it, the complete source code for your application must be also available under the conditions of the GPL (GPL Section 2[b]).
- One of several Quantum Leaps commercial licenses, which are designed for customers who wish to retain the proprietary status of their code and therefore cannot use the GNU General Public License. The customers who license Quantum Leaps software under the commercial licenses do not use the software under the GPL and therefore are not subject to any of its terms.

For more information, please visit the licensing section of our website at: www.state-machine.com/licensing.



2 Directories and Files

The code for the QP port to ARM Cortex-M is part of the **standard QP distribution**, which also contains example applications.

The code of the ARM Cortex-M port is organized according to the Application Note: "[QP_Directory_Structure](#)". Specifically, for this port the files are placed in the following directories:

Listing 1: Directories and files pertaining to the ARM Cortex-M QP port included in the standard QP distribution.

```

<qp>/
|
+-include/
| +-qassert.h
| +-qevent.h
| +-qep.h
| +-qf.h
| +-qk.h
| +-qs.h
| +-qs_dummy.h
| +-qqueue.h
| +-qmpool.h
| +-qpset.h
| +-qvanilla.h
|
+-ports/
| +-arm-cortex/
| | +-vanilla/
| | | +-gnu/
| | | | +-dbg/
| | | | | +-libqep_cortex-m3_cs.a - QEP library for Cortex-M3 with Code Sourcery
| | | | | +-libqf_cortex-m3_cs.a - QF library for Cortex-M3 with Code Sourcery
| | | | +-rel/
| | | | | +-...
| | | | +-spy/
| | | | | +-libqep_cortex-m3_cs.a - QEP library for Cortex-M3 with Code Sourcery
| | | | | +-libqf_cortex-m3_cs.a - QF library for Cortex-M3 with Code Sourcery
| | | | | +-libqs_cortex-m3_cs.a - QS library for Cortex-M3 with Code Sourcery
| | | | +-make_cortex-m3_cs.bat - Batch to build QP for Cortex-M3 Code Sourcery
| | | +-qep_port.h
| | | +-qf_port.h
| | | +-qs_port.h
| | | +-qp_port.h
| | +-qk/
| | | +-gnu/
| | | | +-dbg/
| | | | | +-libqep_cortex-m3_cs.a - QEP library for Cortex-M3 with Code Sourcery
| | | | | +-libqf_cortex-m3_cs.a - QF library for Cortex-M3 with Code Sourcery
| | | | | +-libqk_cortex-m3_cs.a - QK library for Cortex-M3 with Code Sourcery
| | | +-rel/
| | | +-src/
| | | | +-qk_port.s
| | | +-make_cortex-m3_cs.bat
| | | +-qep_port.h

```

```
| | | | +-qf_port.h      - QF platform-dependent public include
| | | | +-qs_port.h      - QS platform-dependent public include
| | | | +-qp_port.h      - QP platform-dependent public include
|
+-examples/              - subdirectory containing the QP example files
| +-arm-cortex/           - ARM Cortex-M port
| | +-vanilla/            - "vanilla" examples (non-preemptive scheduler of QF)
| | | +-gnu/              - GNU ARM compiler
| | | | +-dpp-ev-lm3s811/ - Dining Philosophers example for EV-LM3S811
| | | | | +-cmsis/         - directory containing the CMSIS files
| | | | | +-dbg/          - directory containing the Debug build
| | | | | +-rel/          - directory containing the Release build
| | | | | +-spy/          - directory containing the Spy build
| | | | |
| | | | | +-cproject       - Eclipse project file for the Sourcery CodeBench IDE
| | | | | +-project       - Eclipse project file for the Sourcery CodeBench IDE
| | | | | +-Makefile      - external Makefile for the Sourcery CodeBench IDE
| | | | | +-lm3s811.ld    - linker command file for LM3S811
| | | | | +-lm3s_config.h - CMSIS-compliant configuration for LM3Sxx MCUs
| | | | | +-bsp.c         - Board Support Package for the DPP application
| | | | | +-bsp.h         - BSP header file
| | | | | +-dpp.h         - the DPP header file
| | | | | +-main.c        - the main function
| | | | | +-philos.c      - the Philosopher active object
| | | | | +-table.c       - the Table active object
| | | | |
| | | | | +-game-ev-lm3s811/ - Fly 'n' Shoot game for EV-LM3S811
| | | | | | +-cmsis/       - directory containing the CMSIS files
| | | | | | +-dbg/        - directory containing the Debug build
| | | | | | +-rel/        - directory containing the Release build
| | | | | | +-spy/        - directory containing the Spy build
| | | | | |
| | | | | | +-cproject     - Eclipse project file for the Sourcery CodeBench IDE
| | | | | | +-project     - Eclipse project file for the Sourcery CodeBench IDE
| | | | | | +-Makefile    - external Makefile for the Sourcery CodeBench IDE
| | | | | | +-lm3s811.ld  - linker command file for LM3S811
| | | | | | +-lm3s_config.h - CMSIS-compliant configuration for LM3Sxx MCUs
| | | | | | +-bsp.c       - Board Support Package for this application
| | | | | | +-bsp.h       - BSP header file
| | | | | | +-game.h       - the game header file
| | | | | | +-main.c      - the main function
| | | | | | +-missile.c   - the Missile active object
| | | | | | +-ship.c      - the Ship active object
| | | | | | +-tunnel.c    - the Tunnel active object
| | | | |
| | +-qk/                 - QK examples (preemptive kernel)
| | | +-gnu/              - GNU ARM compiler
| | | | +-dpp-qk-ev-lm3s811/ - DPP example for LPCXpresso-1114
| | | | | |=. . .
| | | | | +-game-qk-ev-lm3s811/ - Fly 'n' Shoot game for EV-LM3S811
| | | | | |=. . .
```

2.1 Building the QP Libraries

All QP components are deployed as libraries that you statically link to your application. The pre-built libraries for QEP, QF, QS, and QK are provided inside the `<qp>\ports\arm-cortex` directory (see [Listing 1](#)). This section describes steps you need to take to rebuild the libraries yourself.

NOTE: To achieve commonality among different development tools, Quantum Leaps software does not use the vendor-specific IDEs, such as the Sourcery CodeBench IDE, for building the QP libraries. Instead, QP supports *command-line* build process based on simple batch scripts.

The code distribution contains the batch file `make_<core>.bat` for building all the libraries located in the `<qp>\ports\arm-cortex\...` directory. For example, to build the debug version of all the QP libraries for ARM Cortex-M, with the GNU ARM compiler, QK kernel, you open a console window on a Windows PC, change directory to `<qp>\ports\arm-cortex\qk\gnu\`, and invoke the batch by typing at the command prompt the following command:

```
make_cortex-m3_cs
```

The build process should produce the QP libraries in the location: `<qp>\ports\arm-cortex\qk\gnu\dbg\`. The batch files assume that the Code Sourcery GNU toolset has been installed in the directory `C:\tools\CodeSourcery\`.

NOTE: You need to adjust the symbol `GNU_ARM` at the top of the batch scripts if you've installed the GNU ARM toolset into a different directory.

In order to take advantage of the QS ("spy") instrumentation, you need to build the QS version of the QP libraries. You achieve this by invoking the `make_cortex-m3_cs.bat` utility with the "spy" target, like this:

```
make_cortex-m3_cs spy
```

The make process should produce the QP libraries in the directory: `<qp>\ports\arm-cortex\vanilla\gnu\spy\`. You choose the build configuration by providing a target to the `make_cortex-m3_cs.bat` utility. The default target is "dbg". Other targets are "rel", and "spy" respectively. The following table summarizes the targets accepted by `make_cortex-m3.bat`.

Table 1 Make targets for the Debug, Release, and Spy software configurations

Software Version	Build commands
Debug (default)	<code>make_cortex-m3_cs</code>
Release	<code>make_cortex-m3_cs rel</code>
Spy	<code>make_cortex-m3_cs spy</code>

2.2 Building and Debugging the Examples

The example applications for ARM Cortex-M have been tested with the EV-LM3S811 evaluation board from Texas Instruments (see [Figure 1](#)) and the GNU/Eclipse-based Code Sourcery CodeBench IDE. The examples contain the Makefile-based Eclipse projects for the CodeBench IDE as well as the Makefiles, so that you can conveniently build and debug the examples both from the CodeBench IDE and from the

command prompt. The provided Makefiles and projects support building the Debug, Release, and Spy configurations.

NOTE: The provided Make files for building the QP applications assume that the GNU ARM toolchain has been installed in the directory `C:\tools\CodeSourcery\`. You need to adjust the symbol `GNU_ARM` at the top of the Makefile to the location of the LPCXpresso installation directory on your system. Alternatively, you can define the `GNU_ARM` symbol as an environment variable, in which case you don't need to modify the Makefile.

NOTE: The provided Make files also assume that you have defined the environment variable `QPC`, if you are using the QP/C framework or the environment variable `QPCPP`, if you are using the QP/C++ framework. These environment variables must contain the paths to the installation directories of the QP/C and QP/C++ frameworks, respectively.

Defining the QP framework locations in environment variables allows you to locate your application in any directory or file system, regardless of the relative path to the QP frameworks.

2.2.1 Building the Examples from Command Line

The example directory `<qp>\examples\arm-cortex\vanilla\gnu\dpp-ev-lm3s811\` contains the Makefile you can use to build the application. The Makefile supports three build configurations: Debug (default), Release, and Spy. You choose the build configuration by defining the `CONF` symbol at the command line, as shown in the table below.

Table 2 Make targets for the Debug, Release, and Spy software configurations

Build Configuration	Build command
Debug (default)	<code>make</code>
Release	<code>make CONF=rel</code>
Spy	<code>make CONF=spy</code>
Clean the Debug configuration	<code>make clean</code>
Clean the Release configuration	<code>make CONF=rel clean</code>
Clean the Spy configuration	<code>make CONF=spy clean</code>

2.2.2 Building the Examples from Eclipse

The example code contains the Eclipse projects for building and debugging the DPP examples with the CodeBench IDE from Code Sourcery. The provided Eclipse projects are Makefile-type projects, which use the same Makefiles that you can call from the command line. In fact the Makefiles are specifically designed to allow building all supported configurations from Eclipse.

NOTE: The provided Makefiles allow you to create and configure the build configurations from the Project | Build Configurations | Manage... sub-menu. For the Release and Spy configurations, you should set the make command to `make CONF=rel` and `make CONF=spy`, respectively. The provided Makefile also correctly supports the clean targets, so invoking Project | Clean... menu for any build configuration works as expected.

3 The Vanilla Port

The “vanilla” port shows how to use QP™ on a “bare metal” ARM Cortex-M based system with the cooperative “vanilla” kernel. In the “vanilla” version of the QP, the only component requiring platform-specific porting is the QF component. The other two components: QEP and QS require merely recompilation and will not be discussed here. With the vanilla port you’re not using the QK component.

3.1 The qep_port.h Header File

The QEP header file for the ARM Cortex-M port is located in `<qp>\ports\arm-cortex\vanilla\gnu\qep_port.h`. Listing 2 shows the `qep_port.h` header file for ARM Cortex-M/GNU. The GNU compiler is a standard C99 compiler, so it simply includes the `<stdint.h>` header file that defines the platform-specific exact-width integer types.

Listing 2: The `qep_port.h` header file for ARM Cortex-M/GNU.

```
#include <stdint.h>                /* C99-standard exact-width integer types */
#include "qep.h"                   /* QEP platform-independent public interface */
```

3.2 The QF Port Header File

The QF header file for the ARM Cortex-M port is located in `<qp>\ports\arm-cortex\vanilla\gnu\qf_port.h`. This file specifies the interrupt locking/unlocking policy (QF critical section) as well as the configuration constants for QF (see Chapter 8 in [PSiCC2]).

The most important porting decision you need to make in the `qf_port.h` header file is the policy for locking and unlocking interrupts. The ARM Cortex-M allows using the simplest “unconditional interrupt unlocking” policy (see Section 7.3.2 of the book “Practical UML Statecharts in C/C++, Second Edition” [PSiCC2]), because ARM Cortex-M is equipped with the standard nested vectored interrupt controller (NVIC) and generally runs ISRs with interrupts unlocked. Listing 3 shows the `qf_port.h` header file for ARM Cortex-M/GNU.

Listing 3: The `qf_port.h` header file for ARM Cortex-M/GNU.

```
/* The maximum number of active objects in the application */
(1) #define QF_MAX_ACTIVE                63

/* The maximum number of event pools in the application */
(2) #define QF_MAX_EPOOL                6

/* QF interrupt disable/enable */
(3) #define QF_INT_DISABLE()             __asm volatile ("cpsid i")
(4) #define QF_INT_ENABLE()             __asm volatile ("cpsie i")

/* QF critical section entry/exit */
(5) /* QF_CRIT_STAT_TYPE not defined: unconditional interrupt unlocking" policy */
(6) #define QF_CRIT_ENTRY(dummy)         __asm volatile ("cpsid i")
(7) #define QF_CRIT_EXIT(dummy)         __asm volatile ("cpsie i")

(8) #include "qep_port.h"                /* QEP port */
(9) #include "qvanilla.h"                /* "Vanilla" cooperative kernel */
(10) #include "qf.h"                     /* QF platform-independent public interface */
```

- (1) The `QF_MAX_ACTIVE` specifies the maximum number of active object priorities in the application. You always need to provide this constant. Here, `QF_MAX_ACTIVE` is set to the maximum limit of 63 active object priorities in the system. You can reduce this number to reduce the RAM footprint of the QF framework.

NOTE: The `qf_port.h` header file does not change the default settings for all the rest of various object sizes inside QF. Please refer to Chapter 8 of [PSiCC2] for discussion of all configurable QF parameters.

- (2) The `QF_MAX_EPOOL` specifies the maximum number of event pools. You don't need to specify this limit, in which case the default of three pools will be chosen.
- (3) The interrupt disable macro resolves to the intrinsic IAR function `__disable_interrupt()`, which in turn generates the single "CPSD i" Thumb2 instruction.
- (4) The unconditional interrupt enable macro resolves to the intrinsic IAR function `__enable_interrupt()`, which in turn generates the single "CPSE i" Thumb2 instruction.
- (5) The `QF_CRIT_STAT_TYPE` is not defined, which means that the simple policy of "unconditional interrupt locking and unlocking" is applied.
- (6) The critical section entry macro resolves to the intrinsic IAR function `__disable_interrupt()`, which in turn generates the single "CPSD i" Thumb2 instruction.
- (7) The critical section exit macro resolves to the intrinsic IAR function `__enable_interrupt()`, which in turn generates the single "CPSE i" Thumb2 instruction.
- (8) This QF port uses the QEP event processor for implementing active object state machines.
- (9) This QF port uses the cooperative "vanilla" kernel.
- (10) The QF port must always include the platform-independent `qf.h` header file.

3.3 Handling Interrupts in the Non-Preemptive Vanilla Kernel

Even though ARM Cortex is designed to use regular C functions as exception and interrupt handlers, in the GNU toolchain functions that are used directly as interrupt handlers should be annotated with `__attribute__((__interrupt__))`. This tells the GNU compiler to add special stack alignment code to the function prologue.

NOTE: Because of a discrepancy between the ARMv7M Architecture and the ARM EABI, it is not safe to use normal C functions directly as interrupt handlers. The EABI requires the stack be 8-byte aligned, whereas ARMv7M only guarantees 4-byte alignment when calling an interrupt vector. This can cause subtle runtime failures, usually when 8-byte types are used [CodeSourcery].

Typically, ISRs are not part of the generic QP port, because it's much more convenient to define ISRs at the application level. The following listing shows all the ISRs in the DPP example application. Please note that the `SysTick_Handler()` ISR calls the `QF_tick()` to perform QF time-event management. (The `SysTick_Handler()` updates also the timestamp used in the QS software tracing instrumentation, see the upcoming Section 8).

```
void SysTick_Handler(void) __attribute__((__interrupt__));
void SysTick_Handler(void) {
#ifdef Q_SPY
    uint32_t dummy = HWREG(NVIC_ST_CTRL); /* clear NVIC_ST_CTRL_COUNT flag */
    QS_tickTime_ += QS_tickPeriod_;        /* account for the clock rollover */
#endif
}
```

```
#endif
    QF_tick();                               /* process all armed time events */
}
```

NOTE: This Application Note complies with the CMSIS standard, which dictates the names of all exception handlers and IRQ handlers.

3.3.1 The Interrupt Vector Table

ARM Cortex-M contains an interrupt vector table (also called the exception vector table) starting usually at address 0x00000000, typically in ROM. The vector table contains the initialization value for the main stack pointer on reset, and the entry point addresses for all exception handlers. The exception number defines the order of entries in the vector table.

ARM Cortex-M requires you to place the initial Main Stack pointer and the addresses of all exception handlers and ISRs into the Interrupt Vector Table allocated typically in ROM. In the GNU compiler, the IDT is initialized in the `startup_LPC11.c` C-language module located in the CMSIS directory.

Listing 4: The interrupt vector table defined in `startup_lm3s.c` (GNU compiler).

```
. . .
/* exception and interrupt vector table -----*/
typedef void (*ExceptionHandler)(void);
typedef union {
    ExceptionHandler handler;
    void *pointer;
} VectorTableEntry;

/* top of stack defined in the linker script */
(1) extern unsigned __c_stack_top__;

/*.....*/
(2) __attribute__((section(".isr_vector")))
VectorTableEntry const g_pfnVectors[] = {
(3)     { .pointer = &__c_stack_top__,          }, /* initial stack pointer */
(4)     { .handler = &Reset_Handler,            }, /* Reset Handler */
        { .handler = &NMI_Handler,              }, /* NMI Handler */
        { .handler = &HardFault_Handler,        }, /* Hard Fault Handler */
        { .handler = &MemManage_Handler,        }, /* MPU Fault Handler */
        { .handler = &BusFault_Handler,         }, /* Bus Fault Handler */
        { .handler = &UsageFault_Handler,       }, /* Usage Fault Handler */
        { .handler = &Spurious_Handler,         }, /* Reserved */
        { .handler = &Spurious_Handler,         }, /* Reserved */
        { .handler = &Spurious_Handler,         }, /* Reserved */
        { .handler = &Spurious_Handler,         }, /* Reserved */
        { .handler = &SVC_Handler,              }, /* SVC Call Handler */
        { .handler = &DebugMon_Handler,        }, /* Debug Monitor Handler */
        { .handler = &Spurious_Handler,         }, /* Reserved */
        { .handler = &PendSV_Handler,          }, /* PendSV Handler */
        { .handler = &SysTick_Handler,         }, /* SysTick Handler */

        /* external interrupts (IRQs) ... */
(5)     { .handler = GPIOPortA_IRQHandler,      }, /* GPIO Port A */
    . . .
```



```
};
```

- (1) The main stack is allocated in its own `.stack` section in the `lm3s811.ld` linker script.

NOTE: The linker script is the place where you determine the stack size. You need to adjust the size of this section to suit your specific application.

NOTE: All QP ports, including the Vanilla port and the QK port use only the main stack (the C-stack). User stack pointer is not used at all.

- (2) The vector table is explicitly placed in the `.isr_vector` section, which the linker script locates at the beginning of Flash at address `0x00000000`.
- (3) The first entry in the IDT is the top of the main stack. The symbol `__c_stack_top` is provided by the linker script.
- (4) The subsequent entries in the IDT are exception handlers. The `Reset_Handler` is defined later in the startup file.
- (5) The third part of the IDT is for interrupt handlers supported by the specific MCU.

3.3.2 Starting Interrupts in `QF_onStartup()`

ARM Cortex-M provides the `SysTick` facility, specifically designed to provide the periodic system time tick, which is configured to deliver the system tick at the `BSP_TICKS_PER_SEC` rate. The configuration of the `SysTick` is done in the `QF_onStartup()` callback using the CMSIS library function `SysTick_Config()`.

QP invokes the `QF_onStartup()` callback just before starting the event loop inside `QF_run()`. The `QF_onStartup()` function is located in the file `bsp.c` and must start the interrupts, in particular the time-tick interrupt.

```
enum ISR_Priorities {    /* ISR priorities starting from the highest urgency */
    GPIOPORTA_PRIO,
    SYSTICK_PRIO,
    /* ... */
};

void QF_onStartup(void) {
    /* set up the SysTick timer to fire at BSP_TICKS_PER_SEC rate */
    SysTick_Config(SystemFrequency / BSP_TICKS_PER_SEC);

    /* set priorities of all interrupts in the system... */
    NVIC_SetPriority(SysTick_IRQn, SYSTICK_PRIO);
    NVIC_SetPriority(GPIOPortA_IRQn, GPIOPORTA_PRIO);

    NVIC_EnableIRQ(GPIOPortA_IRQn);
}
```

3.4 Idle Loop Customization in the “Vanilla” Port

As described in Chapter 7 of [PSiCC2], the “vanilla” port uses the non-preemptive scheduler built into QF. If no events are available, the non-preemptive scheduler invokes the platform-specific callback function `QF_onIdle()`, which you can use to save CPU power, or perform any other “idle” processing (such as Quantum Spy software trace output).

NOTE: The idle callback `QF_onIdle()` must be invoked with interrupts disabled, because the idle condition can be changed by any interrupt that posts events to event queues. `QF_onIdle()` **must** internally enable interrupts, ideally atomically with putting the CPU to the power-saving mode (see also Chapter 7 in [PSiCC2]).

Because `QF_onIdle()` must enable interrupts internally, the signature of the function depends on the interrupt locking policy. In case of the simple “unconditional interrupt locking and unlocking” policy, which is used in this ARM Cortex-M port, the `QF_onIdle()` takes no parameters.

Listing 5 shows an example implementation of `QF_onIdle()` for the LM3S811 MCU. Other ARM Cortex-M embedded microcontrollers (e.g., NXP’s LPC1114/1343) handle the power-saving mode very similarly.

Listing 5: QF_onIdle() callback.

```
(1) void QF_onIdle(void) {           /* entered with interrupts LOCKED, see NOTE01 */
    . . .
(2) #if defined NDEBUG
    /* Put the CPU and peripherals to the low-power mode.
     * you might need to customize the clock management for your application,
     * see the datasheet for your particular ARM Cortex-M MCU.
     */
(3)     __WFI();                      /* Wait-For-Interrupt */
(4)     QF_INT_ENABLE();
    #else
(5)     QF_INT_ENABLE();
    #endif
}
```

- (1) The cooperative Vanilla kernel calls the `QF_onIdle()` callback with interrupts locked, to avoid race condition with interrupts that can post events to active objects and thus invalidate the idle condition.
- (2) The sleep mode is used only in the non-debug configuration, because sleep mode stops CPU clock, which can interfere with debugging.
- (3) The Thumb2 instruction `WFI` (Wait for Interrupt) stops the CPU clock. Note that the CPU stops executing at this line and that interrupts are still **disabled**. An active interrupt first starts the CPU clock again, so the CPU starts executing again. Only after unlocking interrupts in line (4) the interrupt that woke the CPU up is serviced.
- (4-5) The `QF_onIdle()` callback must re-enable interrupts in every path through the code.

NOTE: The idle callback `QF_onIdle()` must unlock interrupts in every path through the code.

4 The QK Port

This section describes how to use QP on ARM Cortex-M with the preemptive QK real-time kernel described in Chapter 10 of [PSiCC2]. The benefit is very fast, fully deterministic task-level response and that execution timing of the high-priority tasks (active objects) will be virtually insensitive to any changes in the lower-priority tasks. The downside is bigger RAM requirement for the stack. Additionally, as with any preemptive kernel, you must be very careful to avoid any sharing of resources among concurrently executing active objects, or if you do need to share resources, you need to protect them with the QK priority-ceiling mutex (again see Chapter 10 of [PSiCC2]).

NOTE: The preemptive configuration with QK uses more stack than the non-preemptive “Vanilla” configuration. You need to adjust the size of this stack to be large enough for your application.

4.1 Single-Stack, Preemptive Multitasking on ARM Cortex-M

The ARM Cortex-M architecture provides a rather unorthodox way of implementing preemptive multitasking, which is designed primarily for the traditional real-time kernels that use multiple per-task stacks. This section explains how the run-to-completion preemptive QK kernel works on ARM Cortex-M.

1. The ARM Cortex-M processor executes application code in the Privileged Thread mode, which is exactly the mode entered out of reset. The exceptions (including all interrupts) are always processed in the Privileged Handler mode.
2. QK uses only the Main Stack Pointer (QK is a single stack kernel). The Process Stack Pointer is not used and is not initialized.
3. The QK port uses the PendSV (exception number 14) and the SVCcall (exception number 11) to perform asynchronous preemptions and context switch, respectively (see Chapter 10 in [PSiCC2]). The application code (your code) **must** initialize the Interrupt Vector Table with the addresses of `PendSV_Handler` and `SVCcall_Handler` exception handlers. Additionally, the interrupt table must be initialized with the SysTick handler that calls `QF_tick()`.
4. The application code (your code) **must** call the function `QK_init()` to set the priority of the PendSV exception to the lowest level in the whole system (0xFF), and the priority of SVCcall to the highest in the system (0x00). The function `QK_init()` sets the priorities of exceptions 14 and 11 to the numerical values of 0xFF and 0x00, respectively. The priorities are set with interrupts disabled, but the interrupt lock key is restored upon the function return.

NOTE: The Stellaris ARM Cortex-M silicon supports only 3 most-significant bits of priority, therefore writing 0xFF to a priority register reads back 0xE0.

5. It is strongly recommended that you do **not** assign the lowest priority (0xFF) to any interrupt in your application. With 3 MSB-bits of priority, this leaves the following 7 priority levels for you (listed from the lowest to the highest urgency): 0xC0, 0xA0, 0x80, 0x60, 0x40, 0x20, and 0x00 (the highest priority).
6. Every ISR **must** set the pending flag for the PendSV exception in the NVIC. This is accomplished in the macro `QK_ISR_EXIT()`, which **must** be called just before exiting from all ISRs (see upcoming Section 4.2.1).
7. ARM Cortex-M enters interrupt context without locking interrupts (without setting the PRIMASK bit). Generally, you should not lock interrupts inside ISRs. In particular, the QF services `QF_publish()`, `QF_tick()`, and `QActive_postFIFO()` should be called with interrupts enabled, to avoid nesting of critical sections.


NOTE: If you don't wish an interrupt to be preempted by another interrupt, you can always prioritize that interrupt in the NVIC to a higher level (use a lower numerical value of priority).

8. In ARM Cortex-M the whole prioritization of interrupts, including the PendSV exception, is performed entirely by the NVIC. Because the PendSV has the lowest priority in the system, the NVIC tail-chains to the PendSV exception only after exiting the last nested interrupt.
9. The restoring of the 8 registers comprising the ARM Cortex-M interrupt stack frame in PendSV is wasteful in a single-stack kernel (see Listing 7(3) and (8)), but is necessary to perform full interrupt return from PendSV to signal End-Of-Interrupt to the NVIC.
10. The pushing of the 8 registers comprising the ARM Cortex-M interrupt stack frame upon entry to SVCALL is wasteful in a single-stack kernel (see Figure 3(10) and (12)), but is necessary to perform full interrupt return to the preempted context through the SVCALL's return.

4.1.1 Examples of Various Preemption Scenarios in QK

Figure 3 illustrates several preemption scenarios in QK.

Figure 3: Various preemption scenarios in the QK preemptive kernel for ARM Cortex-M.



- (0) The timeline in Figure 3 begins with the QK executing the idle loop.
- (1) At some point an interrupt occurs and the CPU immediately suspends the idle loop, pushes the interrupt stack frame to the Main Stack and starts executing the ISR.
- (2) The ISR performs its work, and in QK always sets the pending flag for the **PendSV** exception in the NVIC. The priority of the PendSV exception is configured to be the lowest of all exceptions, so the ISR continues executing and PendSV exception remains pending. At the ISR return, the ARM Cortex-M CPU performs tail-chaining to the pending PendSV exception.


- (3) The whole job of the PendSV exception is to synthesize an interrupt stack frame on top of the stack and perform an interrupt return.
- (4) The PC (exception return address) of the synthesized stack frame is set to `QK_schedule()` (more precisely to a thin wrapper around `QK_schedule()`, see Section 4.3), so the PendSV exception returns to the QK scheduler. The scheduler discovers that the Low-priority task is ready to run (the ISR has posted event to this task). The QK scheduler enables interrupts and launches the Low-priority task, which is simply a C-function call in QK. The Low-priority task (active object) starts running. Some time later another interrupt occurs. The Low-priority task is suspended and the CPU pushes the interrupt stack frame to the Main Stack and starts executing the ISR
- (5) The Low-priority ISR runs and sets the pending flag for the PendSV exception in the NVIC. Before the Low-priority ISR completes, it too gets preempted by a High-priority ISR. The CPU pushes another interrupt stack frame and starts executing the High-priority ISR.
- (6) The High-priority ISR again sets the pending flag for the PendSV exception (setting an already set flag is not an error). When the High-priority ISR returns, the NVIC does not tail-chain to the PendSV exception, because a higher-priority ISR than PendSV is still active. The NVIC performs the normal interrupt return to the preempted Low-priority interrupt, which finally completes.
- (7) Upon the exit from the Low-priority ISR, the NVIC performs tail-chaining to the pending PendSV exception
- (8) The PendSV exception synthesizes an interrupt stack frame to return to the QK scheduler.
- (9) The QK scheduler detects that the High-priority task is ready to run and launches the High-priority task (normal C-function call). The High-priority task runs to completion and returns to the scheduler. The scheduler does not find any more higher-priority tasks to execute and needs to return to the preempted task. The only way to restore the interrupted context in ARM Cortex-M is through the interrupt return, but the task is executing outside of the interrupt context (in fact, tasks are executing in the Privileged Thread mode). The task enters the Handler mode by causing the synchronous **SVC**Call exception
- (10) The only job of the SVCCall exception is to discard its own interrupt stack frame and return using the interrupt stack frame that has been on the stack from the moment of task preemption
- (11) The Low-priority task, which has been preempted all that time, resumes and finally runs to completion and returns to the QK scheduler. The scheduler does not find any more tasks to launch and causes the synchronous SVCCall exception
- (12) The SVCCall exception discards its own interrupt stack frame and returns using the interrupt stack frame from the preempted task context
- (13) The QK idle loop resumes and runs till another interrupt preempts it

4.1.2 The Stack Utilization in QK

Charting the stack utilization over time provides another, complementary view of the preemption scenarios depicted in Figure 3. To demonstrate the essential behavior, the irrelevant function calls and other unrelated stack activity will be ignored.

Figure 4 illustrates the stack utilization across the same preemption scenarios depicted in Figure 3. The timeline and labels used in Figure 3 and Figure 4 are identical to enable easy correlating of these two figures.

Figure 4: Stack depth vs. time for various preemption scenarios in QK for ARM Cortex-M



- (0) The timeline in Figure 4 begins with the QK executing the idle loop.
- (1) At some point an interrupt occurs and the CPU immediately suspends the idle loop, pushes the interrupt stack frame to the Main Stack.
- (2) The ISR starts executing. At the ISR return, the ARM Cortex-M CPU performs tail-chaining to the pending PendSV exception
- (3) Note that tail-chaining keeps the interrupt stack frame unchanged. The PendSV exception synthesizes an interrupt stack frame on top of the stack (2a) and performs an interrupt return
- (3a) The PC (exception return address) of the synthesized stack frame is set to `QK_schedule_()` (more precisely to a thin wrapper around `QK_schedule_()`, see Section 4.3), so the PendSV exception returns to the QK scheduler. The scheduler discovers that the Low-priority task is ready to run (the ISR has posted event to this task). The QK scheduler enables interrupts and launches the Low-priority task, which is simply a C-function call in QK
- (4) The Low-priority task (active object) starts running. Some time later another interrupt occurs. The Low-priority task is suspended and the CPU pushes the interrupt stack frame to the Main Stack
- (5) and starts executing the ISR. Before the Low-priority ISR completes, it too gets preempted by a High-priority ISR. The CPU pushes another interrupt stack frame and starts executing the High-priority ISR
- (6) When the High-priority ISR returns, NVIC performs the normal interrupt return to the preempted Low-priority interrupt
- (7) Upon the exit from the Low-priority ISR, the NVIC performs tail-chaining to the pending PendSV exception. Note that tail-chaining keeps the interrupt stack frame unchanged.
- (7a) The PendSV exception synthesizes an interrupt stack frame as though the interrupt occurred at the beginning of the QK scheduler.
- (8) The PendSV exception returns to the QK scheduler.
- (8a) The QK scheduler detects that the High-priority task is ready to run and launches the High-priority task via a normal C-function call.
- (8b) The High-priority task runs to completion and returns to the scheduler.
- (9) The scheduler causes the synchronous SVCcall exception
- (9a) The only job of the SVCcall exception is to discard its own interrupt stack frame

- (10) The SVCcall exception returns using the interrupt stack frame that has been on the stack from the moment of task preemption.
- (10a) The Low-priority task, which has been preempted all that time, resumes and finally runs to completion and returns to the QK scheduler.
- (11) The scheduler does not find any more tasks to launch and causes the SVCcall exception.
- (11a) The SVCcall exception discards its own interrupt stack frame
- (12) The SVCcall exception returns using the interrupt stack frame from the preempted task context.
- (13) The QK idle loop resumes and runs till another interrupt preempts it.

4.2 The QK Port Header File

In the QK port, you use very similar configuration as the “Vanilla” port described earlier. This section describes only the differences, specific to the QK component.

You configure and customize QK through the header file `qk_port.h`, which is located in the QP ports directory `<qp>\ports\arm-cortex\qk\gnu\`. The most important function of `qk_port.h` is specifying interrupt entry and exit.

NOTE: As any **preemptive** kernel, QK needs to be notified about entering the interrupt context and about exiting an interrupt context in order to perform a context switch, if necessary.

Listing 6: qk_porth.h header file

- ```
(1) #define QK_ISR_ENTRY() do { \
(2) __asm volatile ("cpsid i"); \
(3) ++QK_intNest_; \
(4) QF_QS_ISR_ENTRY(QK_intNest_, QK_currPrio_); \
(5) __asm volatile ("cpsie i"); \
 } while (0)

(6) #define QK_ISR_EXIT() do { \
(7) __asm volatile ("cpsid i"); \
(8) QF_QS_ISR_EXIT(QK_intNest_, QK_currPrio_); \
(9) --QK_intNest_; \
(10) *((uint32_t volatile *)0xE000ED04) = 0x10000000; \
(11) __asm volatile ("cpsie i"); \
 } while (0)

(12) #include "qk.h" /* QK platform-independent public interface */
```
- (1) The `QK_ISR_ENTRY()` macro notifies QK about entering an ISR. The macro body is surrounded by the `do {...} while (0)` loop, which is the standard way of grouping instructions without creating a dangling-else or other syntax problems. In ARM Cortex-M, this macro is called with interrupts unlocked, because the ARM Cortex-M hardware does not set the PRIMASK upon interrupt entry.
  - (2) Interrupts are locked at the ARM Cortex-M core level to perform the following actions atomically.
  - (3) The QK interrupt nesting level `QK_intNest_` is incremented to account for entering an ISR. This prevents invoking the QK scheduler from event posting functions (such as `QActive_postFIFO()` or `QF_publish()`) to perform a synchronous preemption.

- (4) The macro `QF_QS_ISR_ENTRY()` contains the QS instrumentation for interrupt entry. This macro generates no code when QS is inactive (i.e., `Q_SPY` is not defined).
- (5) Interrupts are unlocked at the ARM Cortex-M core level to allow interrupt preemptions.
- (6) The `QK_ISR_EXIT()` macro notifies QK about exiting an ISR.
- (7) Interrupts are locked at the ARM Cortex-M core level to perform the following actions atomically.
- (8) The macro `QF_QS_ISR_EXIT()` contains the QS instrumentation for interrupt exit. This macro generates no code when QS is inactive (i.e., `Q_SPY` is not defined).
- (9) The QK interrupt nesting level `QK_intNest_` is decremented to account for exiting an ISR. This balances step (3).
- (10) This write to the `NVIC_INT_CTRL` register sets the pending flag for the `PendSV` exception.

---

**NOTE:** Setting the pending flag for the `PendSV` exception in every ISR is absolutely **critical** for proper operation of QK. It really does not matter at which point during the ISR execution this happens. Here the `PendSV` is pended at the exit from the ISR, but it could as well be pended upon the entry to the ISR, or anywhere in the middle.

---

- (11) Interrupts are unlocked to perform regular exit from the ISR.
- (12) The QK port header file must include the platform-independent QK interface `qk.h`.

#### 4.2.1 The QK Critical Section

The interrupt locking/unlocking policy in the QK port is the same as in the vanilla port. Please refer to the earlier Section 3.2 for the description of the critical section implementation.



### 4.3 QK Platform-Specific Code for ARM Cortex-M

The QK port to ARM Cortex-M requires coding the `PendSV` and `SVCALL` exceptions in assembly. This ARM Cortex-M-specific code is located in the file `<qp>\ports\arm-cortex\qk\gnu\src\qk_port.s`.

**Listing 7: `QK_init()` function for ARM Cortex-M (file `qk_port.s`)**

```
.syntax unified
.thumb

/*****
 *
 * The QK_init function sets the priorities of PendSV and SVCALL exceptions
 * to 0xFF and 0x00, respectively. The function internally disables
 * interrupts, but restores the original interrupt lock before exit.
 *****/

.section .text.QK_init
.global QK_init
.type QK_init, %function
(1) QK_init:
(2) MRS r0, PRIMASK /* store the state of the PRIMASK in r0 */
(3) CPSID i /* disable interrupts (set PRIMASK) */

(4) LDR r1,=0xE000ED18 /* System Handler Priority Register */
(5) LDR r2,[r1,#8] /* load the System 12-15 Priority Register */
(6) MOVS r3,#0xFF
(7) LSLS r3,r3,#16
(8) ORRS r2,r3 /* set PRI_14 (PendSV) to 0xFF */
(9) STR r2,[r1,#8] /* write the System 12-15 Priority Register */
(10) LDR r2,[r1,#4] /* load the System 8-11 Priority Register */
(11) LSLS r3,r3,#8
(12) BICS r2,r3 /* set PRI_11 (SVCALL) to 0x00 */
(13) STR r2,[r1,#4] /* write the System 8-11 Priority Register */

(14) MSR PRIMASK,r0 /* restore the original PRIMASK */
(15) BX lr /* return to the caller */
.size QK_init, . - QK_init
```

- (1) The `QK_init()` function sets the priorities of the `PendSV` exception (number 14) and `SVCALL` exception (number 11) to the lowest level `0xFF`.
- (2) The `PRIMASK` register is stored in `r0`.
- (3) Interrupts are locked by setting the `PRIMASK`.
- (4) The address of the NVIC System Handler Priority Register 0 is loaded into `r1`.
- (5) The contents of the NVIC System Handler Priority Register 2 (note the offset of 8) is loaded into `r2`.
- (6-7) The mask value of `0xFF0000` is synthesized in `r3`.
- (8) The mask is then applied to set the priority byte `PRI_14` to `0xFF` without changing priority bytes in this register.

- (9) The contents of r2 is stored in the NVIC System Handler Priority Register 2 (note the offset of 8).
- (10) The contents of the NVIC System Handler Priority Register 1 (note the offset of 4) is loaded into r2
- (11) The mask value of 0xFF000000 is synthesized in r3.
- (12) The mask is then applied to set the priority byte PRI\_11 to 0x00 without changing priority bytes in this register.
- (13) The contents of r2 is stored in the NVIC System Handler Priority Register 1 (note the offset of 4).
- (14) The original PRIMASK value is restored.
- (15) The function `QK_init` returns to the caller.

**Listing 8: `PendSV_Handler()` function for ARM Cortex-M (file `qk_port.s`)**

```

/*****
* The PendSV_Handler exception handler is used for handling asynchronous
* preemptions in QK. The use of the PendSV exception is the recommended and
* most efficient method for performing context switches with ARM Cortex.
*
* The PendSV exception should have the lowest priority in the whole system
* (0xFF, see QK_init). All other exceptions and interrupts should have higher
* priority. For example, for NVIC with 2 priority bits all interrupts and
* exceptions must have numerical value of priority lower than 0xC0. In this
* case the interrupt priority levels available to your applications are (in
* the order from the lowest urgency to the highest urgency): 0x80, 0x40, 0x00.
*
* Also, *all* ISRs in the QK application must trigger the PendSV exception
* by calling the QK_ISR_EXIT() macro.
*
* Due to tail-chaining and its lowest priority, the PendSV exception will be
* entered immediately after the exit from the last nested interrupt (or
* exception). In QK, this is exactly the time when the QK scheduler needs to
* check for the asynchronous preemptions.
*****/
.section .text.PendSV_Handler
.global PendSV_Handler
.type PendSV_Handler, %function
(1) .type svc_ret, %function /* to ensure that the svc_ret label is THUMB */

(2) PendSV_Handler:
(3) CPSID i /* disable interrupts at processor level */
(4) BL QK_schedPrio_ /* check if we have preemption */
(5) CMP r0,#0 /* is prio == 0 ? */
(6) BNE.N scheduler /* if prio == 0, branch to scheduler */

(7) CPSIE i /* enable interrupts at processor level */
(8) MOVS r0,#0x6
(9) MVNS r0,r0 /* r0=~0x6=0xFFFFFFFF9 */
(10) BX r0 /* exception-return to the task */

(11) scheduler:
(12) MOVS r3,#1
(13) LSLs r3,r3,#24 /* r3:=(1 << 24), set the T bit (new xpsr) */
(14) LDR r2,=QK_sched_ /* address of the QK scheduler (new pc) */
(15) LDR r1,=svc_ret /* return address after the call (new lr) */

```



```

(16) PUSH {r1-r3} /* push xpsr,pc,lr */
(17) SUB sp,sp,#(4*4) /* don't care for r12,r3,r2,r1 */
(18) PUSH {r0} /* push the prio argument (new r0) */
(19) MOVS r0,#0x6 /* */
(20) MVNS r0,r0 /* r0:=~0x6=0xFFFFFFFF9 */
(21) BX r0 /* exception-return to the scheduler */

(22) svc_ret:
(23) CPSIE i /* enable interrupts to allow SVCcall exception */
(24) SVC 0 /* SV exception returns to the preempted task */

 .size PendSV_Handler, . - PendSV_Handler

```

- (1) The `svc_ret` label is declared as a function so that the GNU assembler/linker synthesizes the right (THUMB) address for that label. (see also step (15) and NOTE following that step).
- (2) The `PendSV_Handler` exception is always entered via tail-chaining from the last nested interrupt (see Section 4.1).
- (3) Interrupts are disabled.
- (4) The function `QK_schedPrio_` is called to find the highest-priority task ready to run. The function is designed to be called with interrupt disabled and returns the priority of this task (in `r0`), or zero if the currently preempted task is the highest-priority.
- (5) If the returned priority is non-zero the code branches to the scheduler.
- (6) This near-branch instruction is reached only if the `QK_schedPrio_` has returned zero, meaning that the PendSV handler needs to return directly to the preempted task.
- (7) Interrupts are enabled.
- (8-9) The special exception-return value `0xFFFFFFFF9` is synthesized in `r0` (two instructions are used to make the code compatible with Cortex-M0, which has no barrel shifter). NOTE: the `r0` register is used instead of `lr` because the Cortex-M0 instruction set cannot manipulate the higher-registers (`r9-r15`).
- (10) This BX instruction causes exception-return to the preempted task. (Exception-return pops the 8-register exception stack frame and changes the processor state to the task-level).
- (11) The scheduler label is reached only when the function `QK_schedPrio_` has returned non-zero task priority. This means that the QK scheduler needs to be invoked to call this task and potentially any tasks that nest on it. The call to the QK scheduler must also perform the mode switch to the task-level.
- (12-13) The value  $(1 \ll 24)$  is synthesized in `r3`. This value is going to be stacked and later restored to xPSR register (only the T bit set).
- (14) The address of the QK scheduler function `QK_sched_` is loaded into `r2`. This will be pushed to the stack as the PC register value.
- (15) The address of the `svc_ret` label is loaded into `r1`. This will be pushed to the stack as the `lr` register value.

**NOTE:** The address of the `svc_ret` label must be a THUMB address, that is, the least-significant bit of this address must be set (this address must be **odd** number). This is essential for the correct return of the QK scheduler with setting the THUMB bit in the PSR. Without the LS-bit set, the ARM Cortex-M CPU will clear the T bit in the PSR and cause the Hard Fault. The GNU assembler/linker will synthesize the correct THUMB address of the `svc_ret` label **only** if this label is declared with the `.type svc_ret, %function` attribute (see step (1)).

- (16) Registers `r3`, `r2` and `r1` are pushed onto the stack.
- (17) The stack pointer is adjusted to leave room for 4 registers. The actual stack contents for these registers is irrelevant.
- (18) The original priority returned in `r0` from `QK_schedPrio_` is pushed to the stack. This will be restored to `r0` register value. This operation completes the synthesis of the exception stack frame. After this step the stack looks as follows:

```

Hi memory
 xPSR
 pc (interrupt return address)
 lr
 r12
 r3
 r2
 r1
old SP --> r0
 xPSR == 0x01000000
 PC == QK_sched_
 lr == svc_ret
 r12 don't care
 r3 don't care
 r2 don't care
 r1 don't care
 SP --> r0 == priority returned from QK_schedPrio_()
Low memory

```

- (19-20) The special exception-return value `0xFFFFFFF9` is synthesized in `r0` (two instructions are used to make the code compatible with Cortex-M0, which has no barrel shifter). NOTE: the `r0` register is used instead of `lr` because the Cortex-M0 instruction set cannot manipulate the higher-registers (`r9-r15`).
- (21) PendSV exception returns using the special value of the `r0` register of `0xFFFFFFF9` (return to Privileged Thread mode using the Main Stack pointer). The synthesized stack frame causes actually a function call to `QK_sched_` function in C.

---

**NOTE:** The return from the PendSV exception just executed switches the ARM Cortex-M core to the Privileged Thread mode. The `QK_sched_` function re-enables interrupts before launching any task, so the tasks always run in the Thread mode with interrupts enabled and can be preempted by interrupts of any priority. This means that as long as the QK priority of the preempting task is higher than the preempted task, the asynchronous preemption can keep repeating at higher and higher levels of nesting on the single stack.

---

- (22) The QK scheduler `QK_sched_()` returns to the `svc_ret` label, because this return address is pushed to the stack in step (14). Please note that the address of the `svc_ret` label must be a THUMB address (see also NOTE after step (14)).
- (23) The interrupts are enabled before calling the SVC exception.
- (24) The synchronous SVC exception is called to put the CPU into the exception mode and correctly return to the thread level.





**Listing 9: SVC\_Handler() function for ARM Cortex-M (file qk\_port.s).**

```

/*****
 *
 * The SVC_Handler exception handler is used for returning back to the
 * interrupted context (task or interrupt). The SVC exception should have
 * the lowest priority in the whole system (see QK_init). The SVCcall
 * exception simply removes its own interrupt stack frame from the stack and
 * returns to the preempted task using the interrupt stack frame that must be
 * at the top of the stack.
 *
 *****/
.section .text.SVC_Handler
.global SVC_Handler
.type SVC_Handler, %function
(1) SVC_Handler:
(2) ADD sp, sp, # (8*4) /* remove one interrupt frame from the stack */
(3) BX lr /* return to the preempted task */
.size SVC_Handler, . - SVC_Handler

.end

```

- (1) The job of SVCcall exception is to discard its own stack frame and cause the interrupt return to the preempted context. The stack contents just before returning from SVCcall exception is shown below:

```

Hi memory
 xPSR
 pc (interrupt return address)
 lr
 r12
 r3
 r2
 r1
SP --> r0
 xPSR don't care
 PC don't care
 lr don't care
 r12 don't care
 r3 don't care
 r2 don't care
 r1 don't care
old SP --> r0 don't care

Low memory

```

- (2) The stack pointer is adjusted to un-stack the 8 registers of the interrupt stack frame corresponding to the SVCcall exception itself.
- (3) SVCcall exception returns using the special value of lr of 0xFFFFFFFF9 (return to Privileged Thread mode using the Main Stack pointer).

## 4.4 Setting up and Starting Interrupts in QF\_onStartup()

Setting up interrupts (e.g., SysTick) for the preemptive QK kernel is identical as in the non-preemptive case. Please refer to Section 3.3.2.

## 4.5 Writing ISRs for QK

QK must be informed about entering and exiting every ISR, so that it can perform asynchronous preemptions. You inform the QK kernel about the ISR entry and exit through the macros `QK_ISR_ENTRY()` and `QK_ISR_EXIT()`, respectively. You need to call these macros in every ISR. The following listing shows the ISR the file `<qp>\examples\arm-cortex\qk\gnu\dpp-qk-lpcexpresso-1114\bsp.c`.

```
void SysTick_Handler(void) __attribute__((__interrupt__));
void SysTick_Handler(void) {
 QK_ISR_ENTRY(); /* inform QK about ISR entry */
#ifdef Q_SPY
 {
 uint32_t dummy = SysTick->CTRL; /* clear the COUNTFLAG in SysTick */
 QS_tickTime_ += QS_tickPeriod_; /* account for the clock rollover */
 }
#endif
 QF_tick();
 QK_ISR_EXIT(); /* inform QK about ISR exit */
}
/*.....*/
void PIOINT0_IRQHandler(void) __attribute__((__interrupt__));
void PIOINT0_IRQHandler(void) {
 QK_ISR_ENTRY(); /* inform QK-nano about ISR entry */
 QActive_postFIFO(AO_Table, Q_NEW(QEvent, MAX_PUB_SIG)); /* for testing */
 QK_ISR_EXIT(); /* inform QK-nano about ISR exit */
}
```

## 4.6 QK Idle Processing Customization in QK\_onIdle()

QK can very easily detect the situation when no events are available, in which case QK calls the `QK_onIdle()` callback. You can use `QK_onIdle()` to suspended the CPU to save power, if your CPU supports such a power-saving mode. Please note that `QK_onIdle()` is called repetitively from the event loop whenever the event loop has no more events to process, in which case only an interrupt can provide new events. The `QK_onIdle()` callback is called with interrupts **unlocked** (which is in contrast to the `QF_onIdle()` callback used in the non-preemptive configuration, see Section 3.4).

The Thumb-2 instruction set used exclusively in ARM Cortex-M provides a special instruction `WFI` (Wait-for-Interrupt) for stopping the CPU clock, as described in the “ARMv7-M Reference Manual” [ARM 06a]. The following Listing 10 shows the `QF_onIdle()` callback that puts ARM Cortex-M into the idle power-saving mode.

**Listing 10: QK\_onIdle() for the preemptive QK configuration.**

```
(1) void QK_onIdle(void) {
 /* toggle the User LED on and then off, see NOTE01 */
(2) QF_INT_LOCK(ignore);
(3) GPIOSetValue(LED_PORT, LED_BIT, LED_ON); /* LED on */
(4) GPIOSetValue(LED_PORT, LED_BIT, LED_OFF); /* LED off */
(5) QF_INT_UNLOCK(ignore);

(6) #ifdef Q_SPY
```

```

 . . .
(7) #elif defined NDEBUG /* sleep mode inteferes with debugging */
 /* put the CPU and peripherals to the low-power mode, see NOTE02
 * you might need to customize the clock management for your application,
 * see the datasheet for your particular ARM Cortex-M MCU.
 */
(8) __WFI(); /* Wait-For-Interrupt */
 #endif
 }

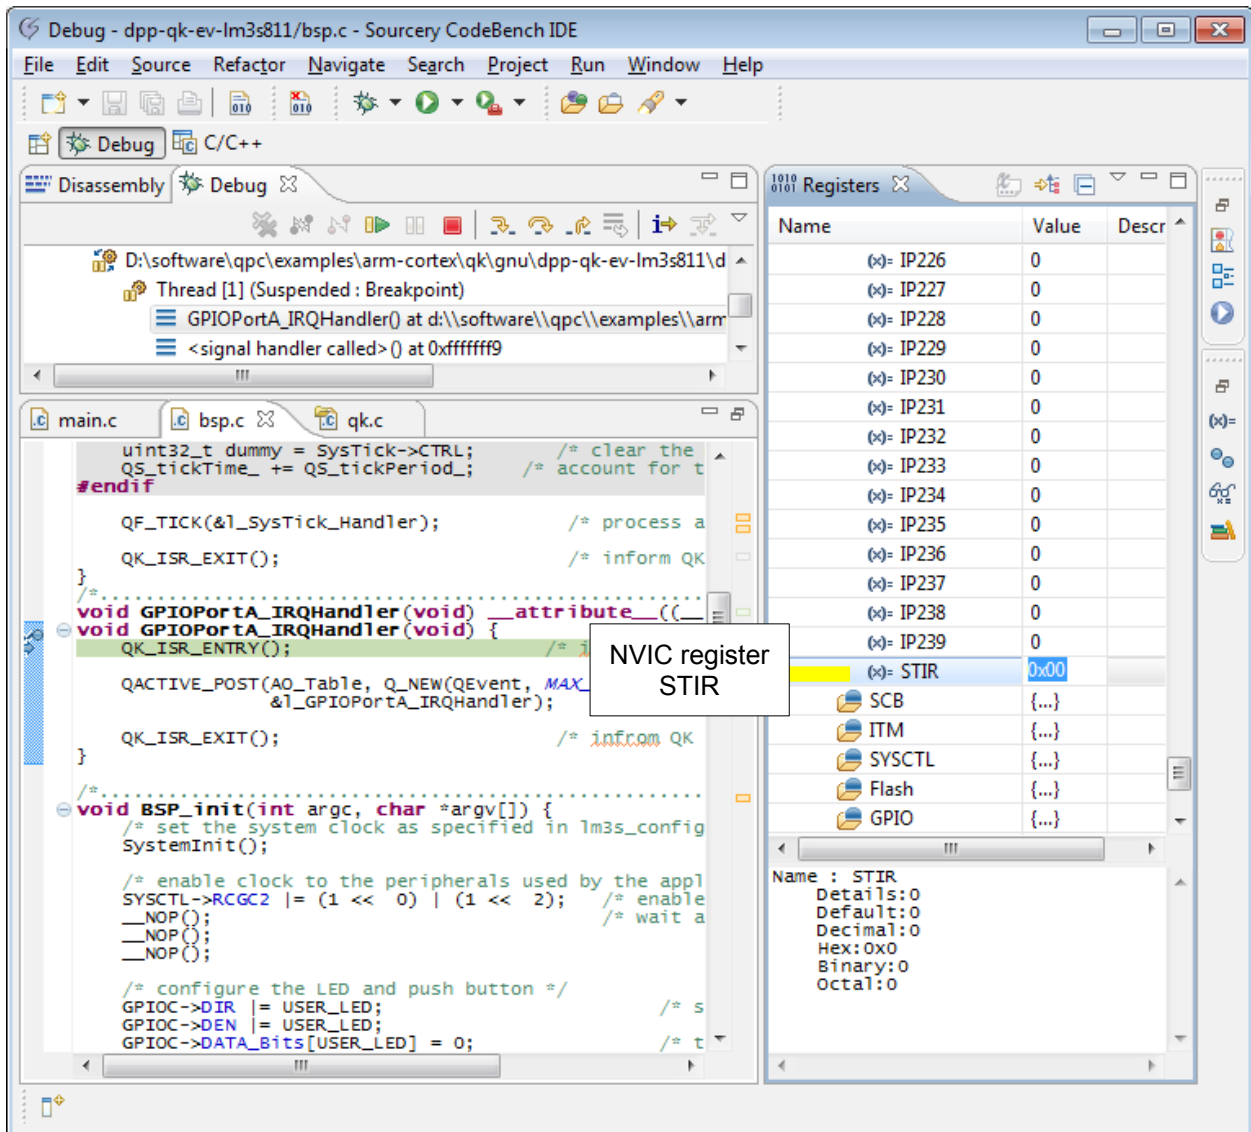
```

- (1) The `QK_onIdle()` function is called with interrupts unlocked.
- (2) The interrupts are locked to prevent preemptions when the LED is on.
- (3-4) Usually, a QK port uses one of the available LEDs to visualize the idle loop activity. However, the LPCXpresso board has only one LED, which is used for other purposes. The code commented out rapidly toggles the LED on and off as long as the idle condition is maintained, so the brightness of the LED is proportional to the CPU idle time (the wasted cycles). Please note that the LED is on in the critical section, so the LED intensity does not reflect any ISR or other processing.
- (5) Interrupts are unlocked.
- (6) This part of the code is only used in the QSpy build configuration. In this case the idle callback is used to transmit the trace data using the UART of the ARM Cortex-M device.
- (7) The following code is only executed when no debugging is necessary (release version).
- (8) The `WFI` instruction is generated using inline assembly.

## 4.7 Testing QK Preemption Scenarios

The DPP example application includes special instrumentation for convenient testing of various preemption scenarios, such as those illustrated in [Figure 5](#).

**Figure 5: Triggering the GPIOA interrupt from the Eclipse debugger.**



The technique described in this section will allow you to trigger an interrupt at any machine instruction and observe the preemptions it causes. The interrupt used for the testing purposes is the GPIOA interrupt (INTID == 0). The ISR for this interrupt is shown below:

```
void GPIOPortA_IRQHandler(void) __attribute__((__interrupt__));
void GPIOPortA_IRQHandler(void) {
 QK_ISR_ENTRY(); /* inform QK-nano about ISR entry */
 QActive_postFIFO(AO_Table, Q_NEW(QEvent, MAX_PUB_SIG)); /* for testing */
 QK_ISR_EXIT(); /* inform QK-nano about ISR exit */
}
```

The ISR, as all interrupts in the system, invokes the macros `QK_ISR_ENTRY()` and `QK_ISR_EXIT()`, and also posts an event to the `Table` active object, which has higher priority than any of the `Philosopher` active object.

[Figure 5](#) shows how to trigger the GPIOA interrupt from the IAR EWARM debugger. From the debugger you need to first open the register window and select NVIC registers from the drop-down list (see right-bottom corner of [Figure 5](#)). You scroll to the STIR register, which denotes the Software Trigger Interrupt Register in the NVIC. This write-only register is useful for software-triggering various interrupts by writing various masks to it. To trigger the GPIOA interrupt you need to write 0x00 to the STIR by clicking on this field, entering the value, and pressing the Enter key.

The general testing strategy is to break into the application at an interesting place for preemption, set breakpoints to verify which path through the code is taken, and trigger the GPIO interrupt. Next, you need to free-run the code (don't use single stepping) so that the NVIC can perform prioritization. You observe the order in which the breakpoints are hit. This procedure will become clearer after a few examples.

#### 4.7.1 Interrupt Nesting Test

The first interesting test is verifying the correct tail-chaining to the `PendSV` exception after the interrupt nesting occurs, as shown in [Figure 3\(7\)](#). To test this scenario, you place a breakpoint inside the `GPIOPortA_IRQHandler()` and also inside the `SysTick_Handler()` ISR. When the breakpoint is hit, you remove the original breakpoint and place another breakpoint at the very next machine instruction (use the Disassembly window) and also another breakpoint on the first instruction of the `QK_PendSV` handler. Next you trigger the `PIOINT0` interrupt per the instructions given in the previous section. You hit the Run button.

The pass criteria of this test are as follows:

1. The first breakpoint hit is the one inside the `GPIOPortA_IRQHandler()` function, which means that GPIO ISR preempted the `SysTick` ISR.
2. The second breakpoint hit is the one in the `SysTick_Handler()`, which means that the `SysTick` ISR continues after the `PIOINT0` ISR completes.
3. The last breakpoint hit is the one in `PendSV_Handler()` exception handler, which means that the `PendSV` exception is tail-chained only after all interrupts are processed.

You need to remove all breakpoints before proceeding to the next test.

#### 4.7.2 Task Preemption Test

The next interesting test is verifying that tasks can preempt each other. You set a breakpoint anywhere in the `Philosopher` state machine code. You run the application until the breakpoint is hit. After this happens, you remove the original breakpoint and place another breakpoint at the very next machine instruction (use the Disassembly window). You also place a breakpoint inside the `GPIOPortA_IRQHandler()` interrupt handler and on the first instruction of the `PendSV_Handler()` handler. Next you trigger the GPIOA interrupt per the instructions given in the previous section. You hit the Run button.

The pass criteria of this test are as follows:



1. The first breakpoint hit is the one inside the `GPIOPortA_IRQHandler()` function, which means that GPIO ISR preempted the `Philosopher` task.
2. The second breakpoint hit is the one in `PendSV_Handler()` exception handler, which means that the PendSV exception is activated before the control returns to the preempted `Philosopher` task.
3. After hitting the breakpoint in QK `PendSV_Handler` handler, you single step into the `QK_scheduler_()`. You verify that the scheduler invokes a state handler from the `Table` state machine. This proves that the `Table` task preempts the `Philosopher` task.
4. After this you free-run the application and verify that the next breakpoint hit is the one inside the `Philosopher` state machine. This validates that the preempted task continues executing only after the preempting task (the `Table` state machine) completes.

#### 4.7.3 Other Tests

Other interesting tests that you can perform include changing priority of the GPIOA interrupt to be lower than the priority of SysTick to verify that the PendSV is still activated only after all interrupts complete.

In yet another test you could post an event to `Philosopher` active object rather than `Table` active object from the `GPIOPortA_IRQHandler()` function to verify that the QK scheduler will not preempt the `Philosopher` task by itself. Rather the next event will be queued and the `Philosopher` task will process the queued event only after completing the current event processing.

## 5 QS Software Tracing Instrumentation

Quantum Spy (QS) is a software tracing facility built into all QP components and also available to the Application code. QS allows you to gain unprecedented visibility into your application by selectively logging almost all interesting events occurring within state machines, the framework, the kernel, and your application code. QS software tracing is minimally intrusive, offers precise time-stamping, sophisticated runtime filtering of events, and good data compression (please refer to “QSP Reference Manual” section in the “QP/C Reference Manual” and also to Chapter 11 in [PSiCC2]).

This QDK demonstrates how to use the QS to generate real-time trace of a running QP application. Normally, the QS instrumentation is inactive and does not add any overhead to your application, but you can turn the instrumentation on by defining the `Q_SPY` macro and recompiling the code.

QS can be configured to send the real-time data out of the serial port of the target device. On the LM3S811 MCU, QS uses the built-in UART to send the trace data out. The EV-LM3S811 board has the UART connected to the virtual COM port provided by the USB debugger (see [Figure 1](#)), so the QSPY host application can conveniently receive the trace data on the host PC. The QS platform-dependent implementation is located in the file `bsp.c` and looks as follows:

**Listing 11: QSpy implementation to send data out of the UART0 of the LM3S811 MCU.**

```
(1) #ifndef Q_SPY

(2) #include "uart.h"

(3) QSTimeCtr QS_tickTime_;
(4) QSTimeCtr QS_tickPeriod_;

(5) enum QSDppRecords {
 QS_PHILO_DISPLAY = QS_USER
 };

 /*.....*/
(6) uint8_t QS_onStartup(void const *arg) {
(7) static uint8_t qsBuf[4*256]; /* buffer for Quantum Spy */
(8) QS_initBuf(qsBuf, sizeof(qsBuf));

 UARTInit(QS_BAUD_RATE); /*initialize the UART with the desired baud rate*/
 NVIC_DisableIRQ(UART_IRQn); /*do not use the interrupts (QS uses polling)*/
 LPC_UART->IER = 0;

 QS_tickPeriod_ = SystemFrequency / BSP_TICKS_PER_SEC;
 QS_tickTime_ = QS_tickPeriod_; /* to start the timestamp at zero */

 return (uint8_t)1; /* return success */
}
/*.....*/
(9) void QS_onCleanup(void) {
}
/*.....*/
(10) void QS_onFlush(void) {
 uint16_t b;
 while ((b = QS_getByte()) != QS_EOD) { /* while not End-Of-Data... */
 while ((LPC_UART->LSR & LSR_THRE) == 0) { /* while TXE not empty */
 }
 LPC_UART->THR = (b & 0xFF); /* put into the THR register */
 }
}
```

```

 }
 }
 /*.....*/
(11) QSTimeCtr QS_onGetTime(void) { /* invoked with interrupts locked */
(12) if ((HWREG(NVIC_ST_CTRL) & NVIC_ST_CTRL_COUNT) == 0) { /* COUNT no set? */
(13) return QS_tickTime_ - (QSTimeCtr)SysTick->VAL;
 }
 else { /* the rollover occurred, but the SysTick_ISR did not run yet */
(14) return QS_tickTime_ + QS_tickPeriod_ - (QSTimeCtr)SysTick->VAL;
 }
 }
}
#endif /* Q_SPY */

```

- (1) The QS instrumentation is enabled only when the macro `Q_SPY` is defined
- (2) The QS implementation uses the UART driver provided in the LPC library.
- (3-4) These variables are used for time-stamping the QS data records. This `QS_tickTime_` variable is used to hold the 32-bit-wide SysTick timestamp at tick. The `QS_tickPeriod_` variable holds the nominal number of hardware clock ticks between two subsequent SysTicks. The SysTick ISR increments `QS_tickTime` by `QS_tickPeriod_`.
- (5) This enumeration defines application-specific QS trace record(s), to demonstrate how to use them.
- (6) You need to define the `QS_init()` callback to initialize the QS software tracing.
- (7) You should adjust the QS buffer size (in bytes) to your particular application
- (8) You always need to call `QS_initBuf()` from `QS_init()` to initialize the trace buffer.
- (9) The `QS_exit()` callback performs the cleanup of QS. Here nothing needs to be done.
- (10) The `QS_flush()` callback flushes the QS trace buffer to the host. Typically, the function busy-waits for the transfer to complete. It is only used in the initialization phase for sending the QS dictionary records to the host (see please refer to “QSP Reference Manual” section in the “QP/C Reference Manual” and also to Chapter 11 in [PSiCC2])

## 5.1 QS Time Stamp Callback `QS_onGetTime()`

The platform-specific QS port must provide function `QS_onGetTime()` (Listing 11(11)) that returns the current time stamp in 32-bit resolution. To provide such a fine-granularity time stamp, the ARM Cortex-M port uses the SysTick facility, which is the same timer already used for generation of the system clock-tick interrupt.

---


**NOTE:** The `QS_onGetTime()` callback is always called with interrupts locked.

---

Figure 6 shows how the SysTick Current Value Register reading is extended to 32 bits. The SysTick Current Value Register (`NVIC_ST_CURRENT`) counts down from the reload value stored in the SysTick Reload Value Register (`NVIC_ST_RELOAD`). When `NVIC_ST_CURRENT` reaches 0, the hardware automatically reloads the `NVIC_ST_CURRENT` counter from `NVIC_ST_RELOAD` on the subsequent clock tick. Simultaneously, the hardware sets the `NVIC_ST_CTRL_COUNT` flag, which “remembers” that the reload has occurred.

The system clock tick ISR `SysTick_Handler()` keeps updating the “tick count” variable `QS_tickTime_` by incrementing it each time by `QS_tickPeriod_`. The clock-tick ISR also clears the `NVIC_ST_CTRL_COUNT` flag.

**Figure 6: Using the SysTick Current Value Register to provide 32-bit QS time stamp.**



Listing 11(11-15) shows the implementation of the function `QS_onGetTime()`, which combines all this information to produce a monotonic time stamp.

- (12) The `QS_onGetTime()` function tests the `NVIC_ST_CTRL_COUNT`. This flag being set means that the `NVIC_ST_CURRENT` has rolled over to zero, but the SysTick ISR has not run yet (because interrupts are still locked).
- (13) Most of the time the `NVIC_ST_CTRL_COUNT` flag is not set, and the time stamp is simply the sum of `QS_tickTime_ + (-HWREG(NVIC_ST_CURRENT))`. Please note that the `NVIC_ST_CURRENT` register is negated to make it to an up-counter rather than down-counter.
- (13) If the `NVIC_ST_CTRL_COUNT` flag is set, the `QS_tickTime_` counter misses one update period and must be additionally incremented by `QS_tickPeriod_`.

## 5.2 QS Trace Output in `QF_onIdle()/QK_onIdle()`

To be minimally intrusive, the actual output of the QS trace data happens when the system has nothing else to do, that is, during the idle processing. The following code snippet shows the code placed either in the `QF_onIdle()` callback ("Vanilla" port), or `QK_onIdle()` callback (in the QK port):

**Listing 12: QS trace output using the UART0 of the Stellaris LM3S811 MCU**

```
#define UART_TXFIFO_DEPTH 16
...
void QK_onIdle(void) {
 ...
 #ifndef Q_SPY
(1) if ((LPC_UART->LSR & LSR_THRE) != 0) { /* is THR empty? */
```

```
uint16_t b;
(2) QF_INT_LOCK(dummy);
(3) b = QS_getByte();
(4) QF_INT_UNLOCK(dummy);
(5) if (b != QS_EOD) { /* not End-Of-Data? */
(6) LPC_UART->THR = (b & 0xFF); /* put into the THR register */
 }
}
#elif defined NDEBUG /* sleep mode interferes with debugging */
 . . .
}
```

- (1) The LSR\_THRE flag is set when the TX FIFO becomes empty.
- (2) Interrupts are locked to call QS\_getByte().
- (3) The function QS\_getByte() returns the byte or QS\_EOD (end-of-data) when no data is available.
- (4) The interrupts are unlocked after the call to QS\_getByte().
- (5) If the data is available
- (6) The byte is inserted into the TX FIFO.

### 5.3 Invoking the QSpy Host Application

The QSPY host application receives the QS trace data, parses it and displays on the host workstation (currently Windows or Linux). For the configuration options chosen in this port, you invoke the QSPY host application as follows (please refer to “QSP Reference Manual” section in the “QP/C Reference Manual” and also to Chapter 11 in [PSiCC2]):

```
qspy -cCOM5 -b115200 -C4
```



## 6 Related Documents and References

### Document

[PSiCC2] "Practical UML Statecharts in C/C++, Second Edition", Miro Samek, Newnes, 2008

[Samek+ 06b] "Build a Super Simple Tasker", Miro Samek and Robert Ward, Embedded Systems Design, July 2006.

[ARMv7-M] "ARM v7-M Architecture Application Level Reference Manual", ARM Limited

[Cortex-M3] "Cortex<sup>™</sup>-M3 Technical Reference Manual", ARM Limited

[CodeSourcery] Sourcery G++ Lite ARM EABI Sourcery G++ Lite Getting Started

[LPC111x] "LPC111x Preliminary user manual Rev. 00.10" — 11 January 2010

[LPC1343] "LPC1311/13/42/43 User manual Rev. 01.01" — 11 January 2010

### Location

Available from most online book retailers, such as [amazon.com](http://amazon.com). See also: <http://www.state-machine.com/psicc2.htm>

<http://www.embedded.com/showArticle.jhtml?articleID=190302110>

Available from <http://infocenter.arm.com/help/>.

Available from <http://infocenter.arm.com/help/>.

<http://www.codesourcery.com/sgpp/lite/arm/portal/doc2861/getting-started.pdf>.

Available in PDF from NXP at:  
<http://ics.nxp.com/support/documents/microcontrollers/pdf/user.manual.lpc1111.lpc1112.lpc1113.lpc1114.pdf>


Available in PDF from NXP at:  
<http://ics.nxp.com/support/documents/microcontrollers/pdf/user.manual.lpc13xx.pdf>

## 7 Contact Information

**Quantum Leaps, LLC**  
103 Cobble Ridge Drive  
Chapel Hill, NC 27516  
USA

+1 866 450 LEAP (toll free, USA only)  
+1 919 869-2998 (FAX)

e-mail: [info@quantum-leaps.com](mailto:info@quantum-leaps.com)  
WEB : <http://www.quantum-leaps.com>  
<http://www.state-machine.com>



*“Practical UML Statecharts in C/C++, Second Edition: Event-Driven Programming for Embedded Systems”,*  
by Miro Samek,  
Newnes, 2008

